

Model_Selection_Training_Cluster_1

```
In [2]: import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import roc_auc_score, accuracy_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
import pickle
```

C:\Users\Dinesh\AppData\Roaming\Python\Python39\site-packages\scipy__init__.py:177: UserWarning: A NumPy version >=1.18.5 and <1.26.0 is required for this version of SciPy (detected version 1.26.4
 warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")


```
In [3]: df = pd.read_csv("Data/Cluster_data.csv")
```

```
In [4]: df.head()
```

```
Out[4]:
```

	age	sex	on_thyroxine	query_on_thyroxine	on_antithyroid_medication	sick	pregnant	t
0	42.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	24.0	0.0	0.0	0.0	0.0	0.0	0.0	
2	47.0	1.0	0.0	0.0	0.0	0.0	0.0	
3	71.0	0.0	1.0	0.0	0.0	0.0	0.0	
4	71.0	0.0	0.0	0.0	0.0	0.0	0.0	

5 rows × 22 columns



```
In [5]: list_of_clusters = df["Cluster"].unique()
```

```
In [6]: list_of_clusters
```

```
Out[6]: array([2, 0, 1], dtype=int64)
```


```
In [7]: cluster_data_2 = df[df["Cluster"] == list_of_clusters[2]]
```

In [8]: `cluster_data_2.head()`

Out[8]:

	age	sex	on_thyroxine	query_on_thyroxine	on_antithyroid_medication	sick	pregnant
40	45.0	1.0	0.0	0.0	0.0	0.0	0.0
88	40.0	0.0	0.0	0.0	0.0	0.0	0.0
89	50.0	0.0	0.0	0.0	0.0	0.0	0.0
91	81.0	1.0	0.0	0.0	0.0	0.0	0.0
116	51.0	1.0	0.0	0.0	0.0	0.0	0.0

5 rows × 22 columns



In [9]: `cluster_feature_2 = cluster_data_2.drop(["Cluster", "Label"], axis = 1)`
`cluster_label_2 = cluster_data_2["Label"]`

In [10]: `x_train, x_test, y_train, y_test = train_test_split(cluster_feature_2, cluster_label_2, test_size=0.3, random_state=42)`

Random Forest

In [11]: `clf = RandomForestClassifier()`

In [12]: `clf.fit(x_train, y_train)`

Out[12]: `RandomForestClassifier()`

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [13]: `clf.score(x_test, y_test)`

Out[13]: 0.9965936739659368

In [14]: `clf.score(x_train, y_train)`

Out[14]: 1.0

In [15]: `param_clf = {"n_estimators": [10, 50, 100, 130, 160, 200, 250],`
`"criteria": ['gini', 'entropy'],`
`"max_depth": range(2, 8, 1),`
`"max_features": ['sqrt', 'log2']`
`}`

```
In [16]: grid = GridSearchCV(estimator=clf, param_grid=param_clf, cv=10, n_jobs=-1,
```

```
In [16]: grid.fit(x_train, y_train)
```

```
Out[16]: GridSearchCV(cv=10, error_score='raise', estimator=RandomForestClassifier(),
                    n_jobs=-1,
                    param_grid={'criterion': ['gini', 'entropy'],
                                'max_depth': range(2, 8),
                                'max_features': ['sqrt', 'log2'],
                                'n_estimators': [10, 50, 100, 130, 160, 200, 250]}))
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [17]: grid.best_params_
```

```
Out[17]: {'criterion': 'entropy',
          'max_depth': 7,
          'max_features': 'sqrt',
          'n_estimators': 10}
```

```
In [17]: clf = RandomForestClassifier(criterion='entropy', max_depth=7, max_features
```

```
In [18]: clf.fit(x_train, y_train)
```

```
Out[18]: RandomForestClassifier(criterion='entropy', max_depth=7, n_estimators=10)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [19]: clf.score(x_test, y_test)
```

```
Out[19]: 0.9961070559610705
```

```
In [20]: clf.score(x_train, y_train)
```

```
Out[20]: 0.9979136240350511
```

```
In [22]: prediction_clf = clf.predict_proba(x_test)
```

```
In [25]: if len(y_test.unique()) == 1:
          clf_score = accuracy_score(y_test, prediction_clf)
        else:
          clf_score = roc_auc_score(y_test, prediction_clf, multi_class="ovr")
```

```
In [26]: clf_score
```

```
Out[26]: 0.9932519073016862
```

KNN

```
In [27]: knn = KNeighborsClassifier()
```

```
In [28]: knn.fit(x_train, y_train)
```

```
Out[28]: KNeighborsClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [29]: knn.score(x_test, y_test)
```

```
Out[29]: 0.9951338199513382
```

```
In [30]: param_knn = {
          'n_neighbors' : [i for i in range(5, 25)],
          'algorithm' : ['auto', 'ball_tree', 'kd_tree', 'brute'],
          'leaf_size' : [10, 15, 20, 25, 30, 35, 40, 50],
          'p' : [1,2],
          'weights' : ['uniform', 'distance']
        }
```

```
In [44]: grid_knn = GridSearchCV(estimator=knn, param_grid=param_knn, cv = 10, error
```

```
In [45]: grid_knn.fit(x_train, y_train)
```

```
Out[45]: GridSearchCV(cv=10, error_score='raise', estimator=KNeighborsClassifier(),
                    n_jobs=-1,
                    param_grid={'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],
                                'leaf_size': [10, 15, 20, 25, 30, 35, 40, 50],
                                'n_neighbors': [5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15,
                                                16, 17, 18, 19, 20, 21, 22, 23, 24],
                                'p': [1, 2], 'weights': ['uniform', 'distance']})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [46]: grid_knn.best_params_
```

```
Out[46]: {'algorithm': 'auto',
          'leaf_size': 10,
          'n_neighbors': 5,
          'p': 1,
          'weights': 'distance'}
```

```
In [31]: knn = KNeighborsClassifier(algorithm='auto', leaf_size=10, n_neighbors = 5,
```

```
In [32]: knn.fit(x_test, y_test)
```

```
Out[32]: KNeighborsClassifier(leaf_size=10, p=1, weights='distance')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [33]: knn.score(x_train, y_train)
```

```
Out[33]: 0.9912372209472147
```

```
In [34]: knn.score(x_test, y_test)
```

```
Out[34]: 1.0
```

```
In [35]: prediction_score = knn.predict_proba(x_test)
```

```
In [36]: if len(y_test.unique()) == 1:
          knn_score = accuracy_score(y_test, prediction_score)
        else:
          knn_score = roc_auc_score(y_test, prediction_score, multi_class="ovr")
```

```
In [37]: knn_score
```

```
Out[37]: 1.0
```

SVM

```
In [39]: svm = SVC()
```

```
In [40]: svm.fit(x_train, y_train)
```

```
Out[40]: SVC()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [41]: svm.score(x_test, y_test)
```

```
Out[41]: 0.9712895377128954
```

```
In [42]: svm.score(x_train, y_train)
```

```
Out[42]: 0.968078447736282
```

```
In [43]: param_svc = {
          'kernel': ['linear', 'poly', 'rbf', 'sigmoid']
        }
```

```
In [63]: grid_svc = GridSearchCV(estimator=svm, param_grid=param_svc, cv = 10, n_jobs=1)
```

```
In [64]: grid_svc.fit(x_train, y_train)
```

```
Out[64]: GridSearchCV(cv=10, error_score='raise', estimator=SVC(), n_jobs=-1,
                      param_grid={'kernel': ['linear', 'poly', 'rbf', 'sigmoid']})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [65]: grid_svc.best_params_
```

```
Out[65]: {'kernel': 'linear'}
```

```
In [44]: svm = SVC(kernel = 'linear')
```

```
In [45]: svm.fit(x_train, y_train)
```

```
Out[45]: SVC(kernel='linear')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [46]: svm.score(x_test, y_test)
```

```
Out[46]: 0.9664233576642336
```

```
In [47]: svm.score(x_train, y_train)
```

```
Out[47]: 0.9609847694554559
```

```
In [58]: prediction_svm = svm.score(x_test, y_test)
```

```
In [59]: prediction_svm
```

```
Out[59]: 0.9664233576642336
```

```
In [ ]:
```

Decision Tree

```
In [60]: dt = DecisionTreeClassifier()
```

```
In [61]: dt.fit(x_train, y_train)
```

```
Out[61]: DecisionTreeClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [62]: dt.score(x_test, y_test)
```

```
Out[62]: 0.9965936739659368
```

```
In [63]: dt.score(x_train, y_train)
```

```
Out[63]: 1.0
```

```
In [64]: param_dt = {"criterion" : ['gini', 'entropy'],
                    "splitter" : ['best', 'random'],
                    "max_depth" : range(2, 40, 1),
                    "min_samples_split" : range(2, 10, 1),
                    "min_samples_leaf" : range(1, 10, 1)
                    }
```

```
In [65]: grid_dt = GridSearchCV(estimator=dt, param_grid=param_dt, n_jobs=3, cv = 10)
```

```
In [66]: grid_dt.fit(x_train, y_train)
```

```
Out[66]: GridSearchCV(cv=10, error_score='raise', estimator=DecisionTreeClassifier(
    ),
          n_jobs=3,
          param_grid={'criterion': ['gini', 'entropy'],
                      'max_depth': range(2, 40),
                      'min_samples_leaf': range(1, 10),
                      'min_samples_split': range(2, 10),
                      'splitter': ['best', 'random']})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [68]: grid_dt.best_params_
```

```
Out[68]: {'criterion': 'gini',
          'max_depth': 33,
          'min_samples_leaf': 1,
          'min_samples_split': 5,
          'splitter': 'random'}
```

```
In [69]: dt = DecisionTreeClassifier(criterion='gini', max_depth=33, min_samples_lea
```

```
In [71]: dt.fit(x_train, y_train)
```

```
Out[71]: DecisionTreeClassifier(max_depth=33, min_samples_split=5, splitter='rando
m')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [72]: dt.score(x_test, y_test)
```

```
Out[72]: 0.9975669099756691
```



```
In [73]: dt.score(x_train, y_train)
```

```
Out[73]: 1.0
```

```
In [74]: prediction_dt = dt.predict_proba(x_test)
```

```
In [76]: dt_score = roc_auc_score(y_test, prediction_dt, multi_class='ovr')
```

```
In [77]: dt_score
```

```
Out[77]: 0.9426115627073657
```

Model Accuracy


```
In [79]: import matplotlib.pyplot as plt
```

In [102]:

```
models = ['Random Forest', 'KNN', 'SVM', 'Decision Tree']

accuracies = [clf_score, knn_score, prediction_svm, dt_score]
colors = ['skyblue', 'salmon', 'lightgreen', 'lightcoral']

plt.figure(figsize=(6, 4))
bars = plt.bar(models, accuracies, color=colors)
plt.xlabel('Models')
plt.ylabel('Accuracy')
plt.title('Accuracy of Different Models')
for bar, acc in zip(bars, accuracies):
    plt.text(bar.get_x() + bar.get_width() / 2, bar.get_height() - 0.55, f'{acc}')
plt.show()
```



Here the KNN is performing well on the cluster data 2 so we are using knn to predict the data

```
In [101]: with open("Models/KNN2/knn.pkl", 'wb') as f:
           pickle.dump(knn, f)
```

In []: